



南開大學
Nankai University

计算机学院
并行程序设计 GPU 作业

NVIIDA Accelerating Applications with CUDA C/C++ 课程

姓名：林语盈
学号：2012174
专业：计算机科学与技术

2022 年 6 月 11 日

摘要

本文介绍了 NVIDIA Accelerating Applications with CUDA C/C++ 课程的主要学习内容并附课程完成截图。

关键字: NVIDIA; CUDA; GPU

目录

1 学习过程和内容	2
1.1 nvcc 基本命令	2
1.2 CUDA 基本编程框架	2
1.3 nsys 性能分析工具	4
1.4 CUDA 程序性能优化	4
1.4.1 SM	4
1.4.2 UM, 数据访存和移动开销	4
1.4.3 asynchronous memory prefetching	4
1.5 Nsight Systems	5
1.6 CUDA streams	5
1.7 手动分配内存	6
1.8 异步内存拷贝	6
A 课程学习截图证明	8
B 向量加法代码	9
B.1 无预取	9
B.2 取至 GPU	10
B.3 将 init 函数在 GPU 端运行	11
B.4 heck 前增加预取	11

1 学习过程和内容

1.1 nvcc 基本命令

查看 GPU 信息

```
1 !nvidia-smi
```

编译并运行

```
1 !nvcc -arch=sm_70 -o name path/hello.cu -run
2 // -run 指定需运行
```

1.2 CUDA 基本编程框架

cpu 端即 host 端, gpu 端即 device 端。通过 host 端, 使函数在 device 端运行并返回。使用 `__global__` 指定 GPU 端的函数, 使用 `GPUFunction(block_num, threads_per_block)` 对 GPU 进行调用。若不加说明, CPU 与 GPU 两端的程序将同时执行, 即可能 GPU 尚未结束时, CPU 端已经结束程序并返回。故, 使用 `cudaDeviceSynchronize()` 进行同步, 使 CPU 端等待 GPU 返回。

`GPUFunction(block_num, threads_per_block)` 的参数指明了分配的 block 数与 threads 数, 总线程数为 `block_num * threads_per_block`。多个线程常常用于将循环展开, 假设循环数为 N, 要处理线程数与 N 不同的情况, 并重新设计循环, 注意每个 block 内的线程数存在上限。此外, 使用 `threadIdx.x` 获得当前的线程编号, 使用 `blockIdx.x` 获得当前的 block 编号, 使用 `blockDim.x` 获得当前 block 维度, 使用 `gridDim.x` 获得当前 grid 维度。实际上, 其内存结构均为 2 维, 可以设计 2 维的数据分配方式。

GPU 拥有独立的存储空间, 使用 `cudaMallocManaged(a, size)` 进行分配, 使用 `cudaFree(a)` 进行释放。

综上, CUDA 基本编程框架如下所示:

CUDA 基本编程框架

```
1 #include <stdio.h>
2 __global__
3 void GPUfunction(float *a, int N)
4 {
5     printf("this is from GPU");
6     int index = threadIdx.x + blockIdx.x * blockDim.x;
7     int stride = blockDim.x * gridDim.x;
8
9     for(int i = index; i < N; i += stride)
10    {
11        //do something with a[i]
12        a[i] = a[i]*2;
13    }
14 }
15
16 void CPUfunction(float *a, int N)
17 {
```

```

18     printf("this is from CPU");
19     for(int i = 0; i < N; ++i)
20     {
21         a[i] = i;
22     }
23 }
24
25 int main()
26 {
27     const int N = 2<<20;
28     size_t size = N * sizeof(float);
29     float *a;
30
31     cudaMallocManaged(&a, size);
32
33     initWith(a, N);
34
35     size_t threadsPerBlock = 256;
36     size_t numberOfBlocks = (N + threadsPerBlock - 1) / threadsPerBlock;
37
38     addVectorsInto<<<numberOfBlocks, threadsPerBlock>>>(a, N);
39
40     cudaDeviceSynchronize();
41
42     //check(a, N);
43     cudaFree(a)
44 }

```

对于每个 cuda 系统函数的调用，都应该进行 error 的处理，可使用如下函数进行。

error 的处理

```

1 #include <stdio.h>
2 #include <assert.h>
3
4 inline cudaError_t checkCuda(cudaError_t result)
5 {
6     if (result != cudaSuccess) {
7         fprintf(stderr, "CUDA Runtime Error: %s\n", cudaGetErrorString(result));
8         assert(result == cudaSuccess);
9     }
10    return result;
11 }
12
13 int main()
14 {
15     ...
16     checkCuda( cudaMallocManaged(&a, size) );
17     ...
18     addVectorsInto<<<numberOfBlocks, threadsPerBlock>>>(...);

```

```
19
20     checkCuda( cudaGetLastError() );
21     checkCuda( cudaDeviceSynchronize() );
22     ...
23     checkCuda( cudaFree(a) );
24
25 }
```

1.3 nsys 性能分析工具

使用 nsys 进行性能分析

```
1 !nsys profile --stats=true ./hello
```

`--stats=true` 指定在屏幕上输出，`./hello` 为待分析的可执行程序。

1.4 CUDA 程序性能优化

1.4.1 SM

SM 即流多处理器，每个 SM 可以处理一个 block，一个 GPU 上拥有一定数量的 SM。性能提升策略：选择恰当的 `grid_size`，使之是给定的 GPU 上 SM 数的倍数。使用如下代码可以获取相关的系统参数。

获取相关的系统参数

```
1 int deviceId;
2 cudaGetDevice(&deviceId); // 获取当前GPU编号
3
4 cudaDeviceProp props;
5 cudaGetDeviceProperties(&props, deviceId); // 获取当前GPU相关参数
6
7 //可以如下设计线程数，N=256*numberOfSMs * 32
8     cudaDeviceGetAttribute(&numberOfSMs, cudaDevAttrMultiProcessorCount, deviceId);
9     int threads_per_block = 256;
10    int number_of_blocks = numberOfSMs * 32;
```

`props` 中的相关参数具体可参见<https://docs.nvidia.com/cuda/cuda-runtime-api/structcudaDeviceProp.html>

1.4.2 UM, 数据访存和移动开销

当数据在 `host` 或 `device` 端被初次访问时，需要进行数据移动，这造成了很大的开销，应当设计函数的运行位置（`host`、`device`）和顺序，尽量减少数据的移动。

1.4.3 asynchronous memory prefetching

优化策略：无论 `cpu` 还是 `gpu` 端，在访问任何数据之前，首先进行异步内存预取。

异步内存预取

```

1 cudaMemPrefetchAsync(pointerToSomeUMData, size, deviceId); // 预取数据至 to
   GPU device, deviceId由上述函数获取
2 cudaMemPrefetchAsync(pointerToSomeUMData, size, cudaCpuDeviceId); // 预取数据 to
   host. cudaCpuDeviceId是内置CUDA变量

```

1.5 Nsight Systems

使用 Nsight Systems 工具加载 nsys profile 生成的.qdrep 文件, 即可方便的将程序各部分的运行细节可视化并进行分析。附录中的几个程序分析截图如图1.1所示。

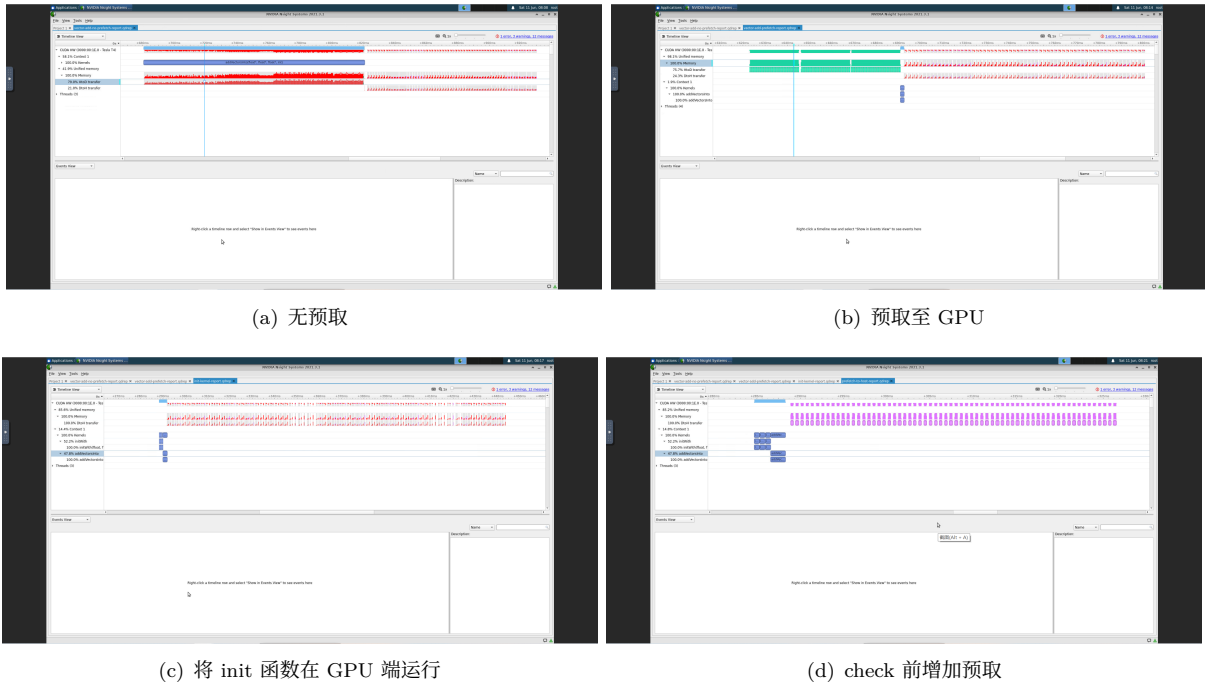


图 1.1: 向量加法程序 Nsight Systems 分析

1.6 CUDA streams

在默认情况下, kernel 函数在 default stream 上运行。但 cuda 提供了对每个 kernel 函数指定特定 stream 的 API。在同一 stream 内, kernel 函数串行运行, 但不同 stream 上的 kernel 函数可以并行运行。default stream 可以进行全部 streams 的同步操作。同样, Nsight Systems 工具可以查看每个 stream 的运行情况。

CUDA streams

```

1 cudaStream_t stream; // 定义 stream
2 cudaStreamCreate(&stream); // 创建 stream
3 GPUfunction<<<1, 1, 0, stream>>>(); // 第四个参数指定了该核函数的 stream
4 // 第三个参数指定分配给每个块的共享内存的字节数, 这里为0
5 cudaStreamDestroy(stream); // 销毁 stream

```

1.7 手动分配内存

CUDA 提供了手动分配 host 和 device 内存的 API, 使用 `cudaMalloc` 与 `cudaMallocHost` 而不再使用 `cudaMallocManaged`, 程序员可以手动在两端分配和回收内存, 并在需要的时候通过 `cudaMemcpy` 进行拷贝。

手动分配内存程序框架

```

1  int *host_a, *device_a;           // Define host-specific and device-specific arrays.
2  cudaMalloc(&device_a, size);      // 'device_a' is immediately available on the GPU.
3  cudaMallocHost(&host_a, size);    // 'host_a' is immediately available on CPU, and is
    page-locked, or pinned.
4
5  initializeOnHost(host_a, N);      // No CPU page faulting since memory is already
    allocated on the host.
6
7  // cudaMemcpy takes the destination, source, size, and a CUDA-provided variable for
    the direction of the copy.
8  cudaMemcpy(device_a, host_a, size, cudaMemcpyHostToDevice);
9
10 kernel<<<blocks, threads, 0, someStream>>>(device_a, N);
11
12 // cudaMemcpy can also copy data from device to host.
13 cudaMemcpy(host_a, device_a, size, cudaMemcpyDeviceToHost);
14
15 verifyOnHost(host_a, N);
16
17 cudaFree(device_a);
18 cudaFreeHost(host_a);             // Free pinned memory like this.

```

1.8 异步内存拷贝

可以将 `cudaMemcpy` 替换为 `cudaMemcpyAsync` 进行异步内存拷贝, 与数据的运算形成流水线, 优化程序性能。但异步时的步长和数据输入应当精心设计, 并分配新的 stream。下例为 4 步内存拷贝。

异步内存拷贝

```

1  /*
2  非异步的内存拷贝:
3  GPUfunction<<<numberOfBlocks, threadsPerBlock>>>(c, a, N);
4  cudaMemcpy(h_c, c, size, cudaMemcpyDeviceToHost);
5  */
6  for (int i = 0; i < 4; ++i)
7  {
8      cudaStream_t stream;
9      cudaStreamCreate(&stream);
10
11     GPUfunction<<<numberOfBlocks/4, threadsPerBlock, 0, stream>>>(&c[i*N/4],
        &a[i*N/4], N/4);

```

```

12     cudaMemcpyAsync(&h_c[i*N/4], &c[i*N/4], size/4, cudaMemcpyDeviceToHost,
13                   stream);
14 }
15 /*
16  const int numberOfSegments = 4;                // This example demonstrates
           slicing the work into 4 segments.
17 int segmentN = N / numberOfSegments;            // A value for a segment's worth
           of N is needed.
18 size_t segmentSize = size / numberOfSegments;    // A value for a segment's worth
           of size is needed.
19
20 // For each of the 4 segments...
21 for (int i = 0; i < numberOfSegments; ++i)
22 {
23     // Calculate the index where this particular segment should operate within the
           larger arrays.
24     segmentOffset = i * segmentN;
25
26     // Create a stream for this segment's worth of copy and work.
27     cudaStream_t stream;
28     cudaStreamCreate(&stream);
29
30     // Asynchronously copy segment's worth of pinned host memory to device over
           non-default stream.
31     cudaMemcpyAsync(&device_array[segmentOffset], // Take care to access correct
           location in array.
32                   &host_array[segmentOffset],    // Take care to access correct
           location in array.
33                   segmentSize,                    // Only copy a segment's worth of
           memory.
34                   cudaMemcpyHostToDevice,
35                   stream);                          // Provide optional argument for
           non-default stream.
36
37     // Execute segment's worth of work over same non-default stream as memory copy.
38     kernel<<<number_of_blocks, threads_per_block, 0,
           stream>>>(&device_array[segmentOffset], segmentN);
39
40     // cudaStreamDestroy will return immediately (is non-blocking), but will not
           actually destroy stream until
41     // all stream operations are complete.
42     cudaStreamDestroy(stream);
43 }
44 */

```


附录 A 课程学习截图证明




图 A.2: 课程学习截图证明 1



图 A.3: 课程学习截图证明 2

附录 B 向量加法代码

B.1 无预取

```
include <stdio.h>
void initWith(float num, float *a, int N) for(int i = 0; i < N; ++i) a[i] = num;

global void addVectorsInto(float *result, float *a, float *b, int N) {
    int index = threadIdx.x + blockIdx.x * blockDim.x;
    int stride = blockDim.x * gridDim.x;
    for(int i = index; i < N; i += stride) result[i] = a[i] + b[i];
}

void checkElementsAre(float target, float *vector, int N) {
    for(int i = 0; i < N; i++) if(vector[i]
```

```

!= target) printf("FAIL: vector[exit(1); printf("Success! All values calculated correctly.");
    int main() int deviceId; int numberOfSMs;
    cudaGetDevice(deviceId); cudaDeviceGetAttribute(numberOfSMs, cudaDevAttrMultiProcessor-
Count, deviceId);
    const int N = 2«24; size_t size = N * sizeof(float);
    float *a; float *b; float *c;
    cudaMallocManaged(a, size); cudaMallocManaged(b, size); cudaMallocManaged(c, size);
    initWith(3, a, N); initWith(4, b, N); initWith(0, c, N);
    size_t threadsPerBlock; size_t numberOfBlocks;
    threadsPerBlock = 256; numberOfBlocks = 32 * numberOfSMs;
    cudaError_t addVectorsErr; cudaError_t asyncErr;
    addVectorsInto«<numberOfBlocks, threadsPerBlock»>(c, a, b, N);
    addVectorsErr = cudaGetLastError(); if(addVectorsErr != cudaSuccess) printf("Error:
    asyncErr = cudaDeviceSynchronize(); if(asyncErr != cudaSuccess) printf("Error:
    checkElementsAre(7, c, N);
    cudaFree(a); cudaFree(b); cudaFree(c);

```

B.2 取至 GPU

```

include <stdio.h>
void initWith(float num, float *a, int N) for(int i = 0; i < N; ++i) a[i] = num;

global void addVectorsInto(float*result, float*a, float*b, int N) int index = threadIdx.x + blockIdx.x * blockDim.x; int stride = blockDim.x * gridDim.x;
for(int i = index; i < N; i += stride) result[i] = a[i] + b[i];
void checkElementsAre(float target, float *vector, int N) for(int i = 0; i < N; i++) if(vector[i]
!= target) printf("FAIL: vector[exit(1); printf("Success! All values calculated correctly.");
    int main() int deviceId; int numberOfSMs;
    cudaGetDevice(deviceId); cudaDeviceGetAttribute(numberOfSMs, cudaDevAttrMultiProcessor-
Count, deviceId);
    const int N = 2«24; size_t size = N * sizeof(float);
    float *a; float *b; float *c;
    cudaMallocManaged(a, size); cudaMallocManaged(b, size); cudaMallocManaged(c, size);
    initWith(3, a, N); initWith(4, b, N); initWith(0, c, N);
    cudaMemPrefetchAsync(a, size, deviceId); cudaMemPrefetchAsync(b, size, deviceId); cudaMemPrefetchA-
sync(c, size, deviceId);
    size_t threadsPerBlock; size_t numberOfBlocks;
    threadsPerBlock = 256; numberOfBlocks = 32 * numberOfSMs;
    cudaError_t addVectorsErr; cudaError_t asyncErr;
    addVectorsInto«<numberOfBlocks, threadsPerBlock»>(c, a, b, N);
    addVectorsErr = cudaGetLastError(); if(addVectorsErr != cudaSuccess) printf("Error:
    asyncErr = cudaDeviceSynchronize(); if(asyncErr != cudaSuccess) printf("Error:
    checkElementsAre(7, c, N);
    cudaFree(a); cudaFree(b); cudaFree(c);

```

B.3 将 init 函数在 GPU 端运行

```

include <stdio.h>

global void initWith(float num, float *a, int N)
int index = threadIdx.x + blockIdx.x * blockDim.x; int stride = blockDim.x * gridDim.x;
for(int i = index; i < N; i += stride) a[i] = num;

global void addVectorsInto(float *result, float *a, float *b, int N)
int index = threadIdx.x + blockIdx.x * blockDim.x; int stride = blockDim.x * gridDim.x;
for(int i = index; i < N; i += stride) result[i] = a[i] + b[i];

void checkElementsAre(float target, float *vector, int N)
for(int i = 0; i < N; i++) if(vector[i] != target) printf("FAIL: vector[");
printf("Success! All values calculated correctly.");

int main()
int deviceId; int numberOfSMs;
cudaGetDevice(deviceId); cudaDeviceGetAttribute(numberOfSMs, cudaDevAttrMultiProcessor-
Count, deviceId);

const int N = 2<<24; size_t size = N * sizeof(float);
float *a; float *b; float *c;
cudaMallocManaged(a, size); cudaMallocManaged(b, size); cudaMallocManaged(c, size);
cudaMemPrefetchAsync(a, size, deviceId); cudaMemPrefetchAsync(b, size, deviceId); cudaMemPrefetchA-
sync(c, size, deviceId);

size_t threadsPerBlock; size_t numberOfBlocks;
threadsPerBlock = 256; numberOfBlocks = 32 * numberOfSMs;
cudaError_t addVectorsErr; cudaError_t asyncErr;
initWith«<numberOfBlocks, threadsPerBlock>>(3, a, N); initWith«<numberOfBlocks, thread-
sPerBlock>>(4, b, N); initWith«<numberOfBlocks, threadsPerBlock>>(0, c, N);
addVectorsInto«<numberOfBlocks, threadsPerBlock>>(c, a, b, N);
addVectorsErr = cudaGetLastError(); if(addVectorsErr != cudaSuccess) printf("Error:
asyncErr = cudaDeviceSynchronize(); if(asyncErr != cudaSuccess) printf("Error:
checkElementsAre(7, c, N);
cudaFree(a); cudaFree(b); cudaFree(c);

```

B.4 heck 前增加预取

```

include <stdio.h>

global void initWith(float num, float *a, int N)
int index = threadIdx.x + blockIdx.x * blockDim.x; int stride = blockDim.x * gridDim.x;
for(int i = index; i < N; i += stride) a[i] = num;

global void addVectorsInto(float *result, float *a, float *b, int N)
int index = threadIdx.x + blockIdx.x * blockDim.x; int stride = blockDim.x * gridDim.x;
for(int i = index; i < N; i += stride) result[i] = a[i] + b[i];

void checkElementsAre(float target, float *vector, int N)
for(int i = 0; i < N; i++) if(vector[i] != target) printf("FAIL: vector[");
printf("Success! All values calculated correctly.");

int main()
int deviceId; int numberOfSMs;
cudaGetDevice(deviceId); cudaDeviceGetAttribute(numberOfSMs, cudaDevAttrMultiProcessor-
Count, deviceId);

const int N = 2<<24; size_t size = N * sizeof(float);

```

```
float *a; float *b; float *c;
cudaMallocManaged(a, size); cudaMallocManaged(b, size); cudaMallocManaged(c, size);
cudaMemPrefetchAsync(a, size, deviceId); cudaMemPrefetchAsync(b, size, deviceId); cudaMemPrefetchA-
sync(c, size, deviceId);
size_t threadsPerBlock; size_t numberOfBlocks;
threadsPerBlock = 256; numberOfBlocks = 32 * numberOfSMs;
cudaError_t addVectorsErr; cudaError_t asyncErr;
initWith«<numberOfBlocks, threadsPerBlock>>(3, a, N); initWith«<numberOfBlocks, thread-
sPerBlock>>(4, b, N); initWith«<numberOfBlocks, threadsPerBlock>>(0, c, N);
addVectorsInto«<numberOfBlocks, threadsPerBlock>>(c, a, b, N);
addVectorsErr = cudaGetLastError(); if(addVectorsErr != cudaSuccess) printf("Error:
asyncErr = cudaDeviceSynchronize(); if(asyncErr != cudaSuccess) printf("Error:
cudaMemPrefetchAsync(c, size, cudaCpuDeviceId);
checkElementsAre(7, c, N);
cudaFree(a); cudaFree(b); cudaFree(c);
```

参考文献